

# Activity\_ Course 7 Salifort Motors project lab

May 31, 2023

## 1 Capstone project: Providing data-driven suggestions for HR

### 1.1 Description and deliverables

This capstone project is an opportunity for you to analyze a dataset and build predictive models that can provide insights to the Human Resources (HR) department of a large consulting firm.

Upon completion, you will have two artifacts that you would be able to present to future employers. One is a brief one-page summary of this project that you would present to external stakeholders as the data professional in Salifort Motors. The other is a complete code notebook provided here. Please consider your prior course work and select one way to achieve this given project question. Either use a regression model or machine learning model to predict whether or not an employee will leave the company. The exemplar following this activity shows both approaches, but you only need to do one.

In your deliverables, you will include the model evaluation (and interpretation if applicable), a data visualization(s) of your choice that is directly related to the question you ask, ethical considerations, and the resources you used to troubleshoot and find answers or solutions.

PACE stages

- Section ??
- Section ??
- Section ??
- Section ??

## 2 Pace: Plan Stage

- Understand your data in the problem context
- Consider how your data will best address the business need
- Contextualize & understand the data and the problem

### Understand the business scenario and problem

The HR department at Salifort Motors wants to take some initiatives to improve employee satisfaction levels at the company. They collected data from employees, but now they don't know what to do with it. They refer to you as a data analytics professional and ask you to provide data-driven suggestions based on your understanding of the data. They have the following question: what's likely to make the employee leave the company?

Your goals in this project are to analyze the data collected by the HR department and to build a model that predicts whether or not an employee will leave the company.

If you can predict employees likely to quit, it might be possible to identify factors that contribute to their leaving. Because it is time-consuming and expensive to find, interview, and hire new employees, increasing employee retention will be beneficial to the company.

### 2.0.1 Familiarize yourself with the HR dataset

The dataset that you'll be using in this lab contains 15,000 rows and 10 columns for the variables listed below.

**Note:** you don't need to download any data to complete this lab. For more information about the data, refer to its source on [Kaggle](#).

Variable	Description
satisfaction_level	Employee-reported job satisfaction level [0–1]
last_evaluation	Score of employee's last performance review [0–1]
number_project	Number of projects employee contributes to
average_monthly_hours	Average number of hours employee worked per month
time_spend_company	How long the employee has been with the company (years)
Work_accident	Whether or not the employee experienced an accident while at work
left	Whether or not the employee left the company
promotion_last_5years	Whether or not the employee was promoted in the last 5 years
Department	The employee's department
salary	The employee's salary (U.S. dollars)

### Reflect on these questions as you complete the plan stage.

- Who are your stakeholders for this project?
- What are you trying to solve or accomplish?
- What are your initial observations when you explore the data?
- What resources do you find yourself using as you complete this stage? (Make sure to include the links.)
- Do you have any ethical considerations in this stage?

```
[20]: ## Relevant Imports for now, will require more later
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
[21]: ## Reading in Data
```

```
df0 = pd.read_csv("HR_capstone_dataset.csv")
df0
```

```
[21]:
```

	satisfaction_level	last_evaluation	number_project	\
0	0.38	0.53	2	
1	0.80	0.86	5	
2	0.11	0.88	7	
3	0.72	0.87	5	
4	0.37	0.52	2	
...	...	...	...	
14994	0.40	0.57	2	
14995	0.37	0.48	2	
14996	0.37	0.53	2	
14997	0.11	0.96	6	
14998	0.37	0.52	2	

	average_monthly_hours	time_spend_company	Work_accident	left	\
0	157	3	0	1	
1	262	6	0	1	
2	272	4	0	1	
3	223	5	0	1	
4	159	3	0	1	
...	...	...	...	...	
14994	151	3	0	1	
14995	160	3	0	1	
14996	143	3	0	1	
14997	280	4	0	1	
14998	158	3	0	1	

	promotion_last_5years	Department	salary
0	0	sales	low
1	0	sales	medium
2	0	sales	medium
3	0	sales	low
4	0	sales	low
...	...	...	...
14994	0	support	low
14995	0	support	low
14996	0	support	low

```
14997          0    support    low
14998          0    support    low
```

```
[14999 rows x 10 columns]
```

```
[22]: ## Exploring Data
```

```
df0.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   satisfaction_level      14999 non-null  float64
1   last_evaluation         14999 non-null  float64
2   number_project          14999 non-null  int64
3   average_monthly_hours  14999 non-null  int64
4   time_spend_company     14999 non-null  int64
5   Work_accident          14999 non-null  int64
6   left                   14999 non-null  int64
7   promotion_last_5years  14999 non-null  int64
8   Department              14999 non-null  object
9   salary                 14999 non-null  object
dtypes: float64(2), int64(6), object(2)
memory usage: 1.1+ MB
```

```
[4]: df0.describe()
```

```
[4]:
```

	satisfaction_level	last_evaluation	number_project \
count	14999.000000	14999.000000	14999.000000
mean	0.612834	0.716102	3.803054
std	0.248631	0.171169	1.232592
min	0.090000	0.360000	2.000000
25%	0.440000	0.560000	3.000000
50%	0.640000	0.720000	4.000000
75%	0.820000	0.870000	5.000000
max	1.000000	1.000000	7.000000

	average_monthly_hours	time_spend_company	Work_accident	left \
count	14999.000000	14999.000000	14999.000000	14999.000000
mean	201.050337	3.498233	0.144610	0.238083
std	49.943099	1.460136	0.351719	0.425924
min	96.000000	2.000000	0.000000	0.000000
25%	156.000000	3.000000	0.000000	0.000000
50%	200.000000	3.000000	0.000000	0.000000
75%	245.000000	4.000000	0.000000	0.000000

max	310.000000	10.000000	1.000000	1.000000
-----	------------	-----------	----------	----------

	promotion_last_5years
count	14999.000000
mean	0.021268
std	0.144281
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

```
[7]: df0.columns
```

```
[7]: Index(['satisfaction_level', 'last_evaluation', 'number_project',
         'average_monthly_hours', 'time_spend_company', 'Work_accident', 'left',
         'promotion_last_5years', 'Department', 'salary'],
        dtype='object')
```

```
[38]: ## Renaming Columns such that they are consistent
```

```
df1 = df0.rename(columns={'Work_accident': 'work_accident', 'Department':
    ↳ 'department', 'average_monthly_hours': 'average_monthly_hours'})
```

```
[39]: ## Checking for null values
```


```
df1.isna().sum()
```

```
[39]: satisfaction_level      0
      last_evaluation        0
      number_project         0
      average_monthly_hours   0
      time_spend_company      0
      work_accident           0
      left                   0
      promotion_last_5years    0
      department             0
      salary                 0
      dtype: int64
```

```
[47]: ## Checking for outliers in time spent at the company
```

```
sns.boxplot(data = df1, y='time_spend_company')
```

```
[47]: <matplotlib.axes._subplots.AxesSubplot at 0x7fba39d47d50>
```



```
[49]: ## Getting rid of outliers
q1 = df1['time_spend_company'].quantile(0.25)
q3 = df1['time_spend_company'].quantile(0.75)
```

```
iqr = q3-q1
print('q1 = ',q1)
print('q3 = ',q3)
print('IQR = ',iqr)
```

```
q1 = 3.0
q3 = 4.0
IQR = 1.0
```

```
[50]: upper_limit = q3 + 1.5*iqr
lower_limit = q1 - 1.5*iqr
print('Lower limit = ',lower_limit)
print('Upper limit = ',upper_limit)
```

```
Lower limit = 1.5
Upper limit = 5.5
```


```
[52]: ## Removing outliers from data set as new data set

df2 = df1.drop(df1[(df1['time_spend_company'] < 1.5) |
↳ (df1['time_spend_company'] > 5.5)].index)
```

```
[54]: ## Displaying Results
```

```
sns.boxplot(data = df2, y='time_spend_company')
```


```
[54]: <matplotlib.axes._subplots.AxesSubplot at 0x7fba37446990>
```



```
[56]: ## Checking for outliers in average monthly hours
```

```
sns.boxplot(data = df2, y='average_monthly_hours')
```

```
[56]: <matplotlib.axes._subplots.AxesSubplot at 0x7fba34f49250>
```



## 2.0.2 Data visualizations

Plotting relevant variables

```
[94]: df2.reset_index(inplace=True,drop=True)
df2
```

```
[94]:
```

	satisfaction_level	last_evaluation	number_project	\
0	0.38	0.53	2	
1	0.11	0.88	7	
2	0.72	0.87	5	
3	0.37	0.52	2	
4	0.41	0.50	2	
...	...	...	...	
13712	0.40	0.57	2	
13713	0.37	0.48	2	
13714	0.37	0.53	2	
13715	0.11	0.96	6	
13716	0.37	0.52	2	

	average_monthly_hours	time_spend_company	work_accident	left	\
0	157	3	0	1	
1	272	4	0	1	
2	223	5	0	1	
3	159	3	0	1	
4	153	3	0	1	



```

...
13712      151      3      0      1
13713      160      3      0      1
13714      143      3      0      1
13715      280      4      0      1
13716      158      3      0      1

```

```

      promotion_last_5years department salary
0              0      sales      low
1              0      sales  medium
2              0      sales      low
3              0      sales      low
4              0      sales      low

```

```

...
13712      0  support      low
13713      0  support      low
13714      0  support      low
13715      0  support      low
13716      0  support      low


```

[13717 rows x 10 columns]

```


[89]: sns.barplot(data=df2, x='time_spend_company', y='satisfaction_level')
      plt.xticks(rotation=0, fontsize=25)
      sns.set(rc={'figure.figsize': (25, 12)})

```




Satisfaction trends downward with time spend at company

```
[90]: sns.boxplot(data=df2, x='promotion_last_5years', y='average_monthly_hours')
plt.xticks(rotation=0, fontsize=25)
sns.set(rc={'figure.figsize': (25, 12)})
```




Average monthly hours is not strongly correlated to whether they've had a promotion in the last 5 years

```
[81]: sns.boxplot(data=df2, x='number_project', y='average_monthly_hours')
plt.xticks(rotation=0, fontsize=20)
sns.set(rc={'figure.figsize': (25, 12)})
```




The more projects an employee contributes to is a good indicator of how many hours they average per month

```
[82]: sns.boxplot(data=df2, x='left', y='average_monthly_hours')
plt.xticks(rotation=0, fontsize=20)
sns.set(rc={'figure.figsize': (25, 12)})
```




The employees that were working the most hours per month on average are more likely to have left

```
[92]: sns.barplot(data=df2, x='salary', y='left')
plt.xticks(rotation=0, fontsize=25)
sns.set(rc={'figure.figsize': (25, 12)})
```




Less likely to leave with higher salary

```
[93]: sns.barplot(data=df2, x='left', y='satisfaction_level')
plt.xticks(rotation=0, fontsize=25)
sns.set(rc={'figure.figsize': (25, 12)})
```



Those that leave have a lower satisfaction level

```
[95]: sns.barplot(data=df2, x='department', y='left')
plt.xticks(rotation=0, fontsize=25)
sns.set(rc={'figure.figsize': (25, 12)})
```




```
[111]: ## Validating barplot

sales = df2[df2['department'] == 'sales']
sales_left = sales.left.sum()
sales_emp = len(sales)
print(sales_left)
print(sales_emp)
percent = sales_left / sales_emp
print(percent)
```

```
966
3739
0.25835784969243114
```

```
[112]: sns.boxplot(data=df2, x='left', y='last_evaluation')
plt.xticks(rotation=0, fontsize=25)
sns.set(rc={'figure.figsize': (25, 12)})
```



Not highly correlated

```
[119]: ## Going to create a heatmap for visualzing correlation matrix of the data, to
        ↪ do this I need all categorical variables to be dummied

df3 = pd.get_dummies(df2, columns=['department', 'salary'])
df3
```

[119]:

	satisfaction_level	last_evaluation	number_project	\
0	0.38	0.53	2	
1	0.11	0.88	7	
2	0.72	0.87	5	
3	0.37	0.52	2	
4	0.41	0.50	2	
...	...	...	...	
13712	0.40	0.57	2	
13713	0.37	0.48	2	
13714	0.37	0.53	2	
13715	0.11	0.96	6	
13716	0.37	0.52	2	

	average_monthly_hours	time_spend_company	work_accident	left	\
0	157	3	0	1	
1	272	4	0	1	
2	223	5	0	1	
3	159	3	0	1	
4	153	3	0	1	
...	...	...	...	...	
13712	151	3	0	1	
13713	160	3	0	1	
13714	143	3	0	1	
13715	280	4	0	1	
13716	158	3	0	1	

	promotion_last_5years	department_IT	department_RandD	...	\
0	0	0	0	...	
1	0	0	0	...	
2	0	0	0	...	
3	0	0	0	...	
4	0	0	0	...	
...	...	...	...	...	
13712	0	0	0	...	
13713	0	0	0	...	
13714	0	0	0	...	
13715	0	0	0	...	
13716	0	0	0	...	

	department_hr	department_management	department_marketing	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	
...	...	...	...	
13712	0	0	0	

13713	0	0	0
13714	0	0	0
13715	0	0	0
13716	0	0	0

	department_product_mng	department_sales	department_support	\
0	0	1	0	
1	0	1	0	
2	0	1	0	
3	0	1	0	
4	0	1	0	
...	...	...	...	
13712	0	0	1	
13713	0	0	1	
13714	0	0	1	
13715	0	0	1	
13716	0	0	1	

	department_technical	salary_high	salary_low	salary_medium
0	0	0	1	0
1	0	0	0	1
2	0	0	1	0
3	0	0	1	0
4	0	0	1	0
...	...	...	...	...
13712	0	0	1	0
13713	0	0	1	0
13714	0	0	1	0
13715	0	0	1	0
13716	0	0	1	0


[13717 rows x 21 columns]

```
[120]: from scipy import stats

fig = df3.corr()

sns.heatmap(fig, cmap='coolwarm', center=0)
```

[120]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7fba20205590>



Whether an employee leaves is most highly correlated to their time spent at the company, whether they've had a work accident, and their salary.

### 2.0.3 Modeling

[139]: *## First going to check the ratio of employees that have left and not,  
## to make sure there is a relatively equal balance*

```
left_df = df3[df3['left'] == 1]
stayed_df = df3[df3['left'] == 0]
left = len(left_df)
stayed = len(stayed_df)
percent = left / (left + stayed)
print(left)
print(stayed)
print(percent * 100)
```

3362

10355

24.509732448786178

25% is very acceptable so we do not need to upsample or downsample

[143]: *## Going to use an XGB Classifier to predict whether an employee is likely to  
↪ leave*



```

from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score,\
f1_score, confusion_matrix, ConfusionMatrixDisplay
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from xgboost import plot_importance

## First need to Identify X and y variables

y = df3['left']

X = df3.drop(columns=['left'], axis = 1)

## Splitting up data into training and testing data
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y,
→test_size=0.2, random_state=42)

```

[142]: y\_test

```

[142]: 11393    1
      4677    0
      10001   0
      13495   1
      8169    0
      ..
      10982   0
      12152   0
      12441   0
      5592    0
      10645   0
      Name: left, Length: 2744, dtype: int64

```

[125]: *## Instantiating Classifier*

```

xgb = XGBClassifier(objective='binary:logistic', random_state=42)

```

[126]: *## Instantiating a GridSearch with hyper parameters*  
*## Refitting with f1 score*

```

cv_params = {'max_depth': [4,8,12],
             'min_child_weight': [3, 5],
             'learning_rate': [0.01, 0.1],
             'n_estimators': [300, 500]
            }

scoring = {'accuracy', 'precision', 'recall', 'f1'}

```

```
xgb_cv = GridSearchCV(xgb, cv_params, scoring=scoring, cv=4, refit='f1')
```

```
[127]: ## Fitting model to training data
```

```
xgb_cv.fit(X_train,y_train)
```

```
[127]: GridSearchCV(cv=4, error_score=nan,  
                  estimator=XGBClassifier(base_score=None, booster=None,  
                                          callbacks=None, colsample_bylevel=None,  
                                          colsample_bynode=None,  
                                          colsample_bytree=None,  
                                          early_stopping_rounds=None,  
                                          enable_categorical=False, eval_metric=None,  
                                          gamma=None, gpu_id=None, grow_policy=None,  
                                          importance_type=None,  
                                          interaction_constraints=None,  
                                          learning_rate=None, max...  
                                          num_parallel_tree=None,  
                                          objective='binary:logistic',  
                                          predictor=None, random_state=42,  
                                          reg_alpha=None, ...),  
                  iid='deprecated', n_jobs=None,  
                  param_grid={'learning_rate': [0.01, 0.1], 'max_depth': [4, 8, 12],  
                              'min_child_weight': [3, 5],  
                              'n_estimators': [300, 500]},  
                  pre_dispatch='2*n_jobs', refit='f1', return_train_score=False,  
                  scoring={'recall', 'accuracy', 'precision', 'f1'}, verbose=0)
```

## 2.0.4 Evaluation

```
[128]: xgb_cv.best_params_
```

```
[128]: {'learning_rate': 0.1,  
        'max_depth': 12,  
        'min_child_weight': 3,  
        'n_estimators': 300}
```

```
[129]: xgb_cv.best_score_
```

```
[129]: 0.9706812495308832
```

```
[131]: ## Evaluating on test data now
```

```
y_pred = xgb_cv.predict(X_test)
```

```
[134]: def get_test_scores(model_name:str, preds, y_test_data):

    accuracy = round(accuracy_score(y_test_data, preds), 3)
    precision = round(precision_score(y_test_data, preds), 3)
    recall = round(recall_score(y_test_data, preds), 3)
    f1 = round(f1_score(y_test_data, preds), 3)

    table = pd.DataFrame({'model': [model_name],
                          'precision': [precision],
                          'recall': [recall],
                          'f1': [f1],
                          'accuracy': [accuracy]
                          })

    return table
```

```
[135]: get_test_scores('XGBClassifier', y_pred, y_test)
```


```
[135]:
```

	model	precision	recall	f1	accuracy
0	XGBClassifier	0.986	0.967	0.977	0.989

```
[137]: ## Plotting confusion matrix to visualse model performance


cm = confusion_matrix(y_test, y_pred, labels=xgb_cv.classes_)

disp = ConfusionMatrixDisplay(confusion_matrix=cm,
                              display_labels=xgb_cv.classes_)
disp.plot();
```



```
[141]: ## Plotting relative importance of each variable used to determine whether  
## Employee is likely to leave or not
```

```
plot_importance(xgb_cv.best_estimator_, max_num_features=10);
```



Model summary: 2100 true negatives, 650 true positives, 22 false negatives, 9 false positives

Overall very strong model, would recommend to use this model to predict whether future employees are likely to leave based on metrics given initially. From the feature importance graph, we can see that employees who worked longer hours are more likely to leave, can be interpreted as overworked. Employees with low satisfaction scores are also more likely to leave. And employees with low evaluation scores are likely to leave.